

# Optimization of GPU Cluster Job Scheduling via Reinforcement Learning

- Based on Alibaba PAI-v2020 Cluster Trace

# Problem Definition & Motivation

- The Problem:
  - Modern AI/ML workloads vary in size, duration, and GPU demand.
  - Standard FIFO scheduling causes resource fragmentation and high queue delays.
- Commercial Value:
  - 1000 GPUs @ \$3/hr = ~\$2.16M/month.
  - 10% efficiency gain = ~\$2.6M/year savings.
  - Objective: Minimize job queuing time & maximize GPU utilization using RL.

# Dataset

- Source: Alibaba Cluster Trace Program (v2020)
- Scale: 6,500+ GPUs, ~1,800 machines, 2 months (Jul–Aug 2020)
- Workload: MLaaS – deep learning training & inference
- Key Features:
  - Heterogeneous GPU types (T4, V100, P100)
- Tables: `pai_job_table`, `pai_machine_spec`,  
`pai_machine_metric`, `pai_group_tag_table`

# Key Challenges & Direction

- Observed Problems:
- Short Task Delays: Many short-duration jobs waited >50% of their runtime in queue.
- High-end GPU Scheduling Issues: Few critical jobs require specific GPUs (V100/NVLink) with topology & concurrency constraints.
- CPU Bottleneck: Data preprocessing (RecSys, GNN, RL) saturates CPU → GPU idle.
- Load Imbalance: Weak-GPU nodes overloaded; high-end GPUs underutilized.
- Resource Mismatch: 2-GPU vs 8-GPU servers have poor CPU/GPU ratio alignment.
  
- Proposed Solutions:
- Reserving-and-Packing Strategy
  - Reserve top GPUs (V100) for demanding jobs, pack smaller tasks onto low-end nodes.
- CPU/GPU Co-optimization
  - Consider multi-resource scheduling (CPU, GPU, Memory) to avoid bottlenecks.

# Data Preprocessing

- Cleaning:
  - - Remove missing plan\_cpu/mem/duration
  - - Keep only successful jobs
- Merging:
  - - pai\_job → pai\_machine
- Result: Unified Job + Machine + Runtime view
- Feature Extraction: state\_cols (cluster load), critical\_cols (job reqs)

# Simulating Heterogeneity

- Challenge: Same job runs faster on V100 than T4
- Method:
  - 1. Group by group\_tag (graphlearn, ctr prediction, bert etc.)
  - 2. Compute avg duration across GPU types
  - 3.  $k = \text{AvgDuration(V100)}/\text{AvgDuration(T4)}$
- Application:  $\text{Predicted\_T4} = \text{Duation\_V100} / k$
- Contribution: Modeling heterogeneous GPU speeds → more realistic simulator.

# RL Agent Design

- Flow:
  - 1. Sort jobs by start\_time
  - 2. Update cluster state
  - 3. RL decides: assign or wait
- State: cluster load, plan\_cpu/mem/gpu
- Action: select machine or wait
- Reward:  $r = -(current\_time - job\_arrival) \& gpu$
- Goal: minimize job waiting delay, maximize gpu utilization

# Baselines & Experimental Setup

- Baselines:
  - FIFO: common but inefficient
  - SJF: favors short jobs, risks starvation
- Simulation:
  - Custom Python simulator with real traces
- Metrics: Avg Waiting Time, GPU Utilization

# Next Steps

- Immediate:
  - Train RL agent vs FIFO/SJF
  - Visualize waiting time & utilization
- Optimization:
  - Tune Hyperparameter, reward scaling
  - Try DQN/PPO
- Goal: Job scheduling + Open-source simulator